

# Lecture 8: Pandas - Hierarchical Indexing

## Data Science, DST, UIC

- Up to this point we've been focused primarily on **one-dimensional and two-dimensional data**, stored in Pandas `Series` and `DataFrame` objects, respectively.
- Often it is useful to go beyond this and store higher-dimensional data—that is, data indexed by **more than one or two keys**.

To handle three-dimensional and four-dimensional data, a common pattern in practice is to make use of **hierarchical indexing** (also known as **multi-indexing**) to incorporate multiple index *levels* within a single index. In this way, higher-dimensional data can be compactly represented within the familiar one-dimensional `Series` and two-dimensional `DataFrame` objects.

In this section, we'll explore the direct creation of `MultiIndex` objects, considerations when indexing, slicing, and computing statistics across multiply indexed data, and useful routines for converting between simple and hierarchically indexed representations of your data.

We begin with the standard imports:

```
In [1]: import numpy as np
import pandas as pd
```

## 1 A Multiply Indexed Series

Let's start by considering how we might represent two-dimensional data within a one-dimensional `Series`. For concreteness, we will consider a series of data where each point has a string and numerical key.

### 1.1 The bad way

Suppose you would like to track data about states from two different years. Using the Pandas tools we've already covered, you might be tempted to simply use Python tuples as keys:

```
In [2]: index = [('California', 2000), ('California', 2010),
                ('New York', 2000), ('New York', 2010),
                ('Texas', 2000), ('Texas', 2010)]
populations = [33871648, 37253956,
               18976457, 19378102,
               20851820, 25145561]
pop = pd.Series(populations, index=index)
pop
```

```
Out[2]: (California, 2000)    33871648
        (California, 2010)    37253956
        (New York, 2000)     18976457
        (New York, 2010)     19378102
        (Texas, 2000)        20851820
        (Texas, 2010)        25145561
dtype: int64
```

With this indexing scheme, you can straightforwardly index or slice the series based on this multiple index:

```
In [3]: pop[('California', 2010):('Texas', 2000)]
```

```
Out[3]: (California, 2010)    37253956
        (New York, 2000)    18976457
        (New York, 2010)    19378102
        (Texas, 2000)       20851820
        dtype: int64
```

But the convenience ends there. For example, if you need to select all values from 2010, you'll need to do some messy (and potentially slow) munging to make it happen:

```
In [4]: pop[[i for i in pop.index if i[1] == 2010]] # [(C, 2010), (NY, 2010), ...]
```

```
Out[4]: (California, 2010)    37253956
        (New York, 2010)     19378102
        (Texas, 2010)        25145561
        dtype: int64
```

This produces the desired result, but is not as clean (or as efficient for large datasets) as the slicing syntax we've grown to love in Pandas.

## 1.2 The Better Way: Pandas MultiIndex

Fortunately, Pandas provides a better way. Our tuple-based indexing is essentially a rudimentary multi-index, and the `Pandas MultiIndex` type gives us the type of operations we wish to have. We can create a multi-index from the tuples as follows:

```
In [5]: index = pd.MultiIndex.from_tuples(index)
        index
```

```
Out[5]: MultiIndex([('California', 2000),
                    ('California', 2010),
                    ('New York', 2000),
                    ('New York', 2010),
                    ('Texas', 2000),
                    ('Texas', 2010)],
                  )
```

You can think of `MultiIndex` as an array of tuples where each tuple is unique.

If we re-index our series with this `MultiIndex`, we see the hierarchical representation of the data:

```
In [6]: pop = pop.reindex(index)
        pop
```

```
Out[6]: California 2000    33871648
          2010    37253956
          New York 2000    18976457
          2010    19378102
          Texas   2000    20851820
          2010    25145561
          dtype: int64
```

Here the first two columns of the `Series` representation show the multiple index values, while the third column shows the data. Notice that some entries are missing in the first column: in this multi-index representation, any blank entry indicates the same value as the line above it.

Now to access all data for which the second index is 2010, we can simply use the Pandas slicing notation:

```
In [7]: pop[:, 2010]

Out[7]: California    37253956
        New York      19378102
        Texas         25145561
        dtype: int64
```

The result is a singly indexed array with just the keys we're interested in. This syntax is much more convenient (and the operation is much more efficient!) than the tuple-based multi-indexing solution that we started with. We'll now further discuss this sort of indexing operation on hierarchically indexed data.

### 1.3 MultiIndex as extra dimension

You might notice something else here: we could easily have stored the same data using a simple `DataFrame` with index and column labels. In fact, Pandas is built with this equivalence in mind. The `unstack()` method will quickly convert a multiply indexed `Series` into a conventionally indexed `DataFrame`:

```
In [8]: pop_df = pop.unstack()
        pop_df
        # unstack the levels(default -1) from index to columns
```

```
Out[8]:
```

|            | 2000     | 2010     |
|------------|----------|----------|
| California | 33871648 | 37253956 |
| New York   | 18976457 | 19378102 |
| Texas      | 20851820 | 25145561 |

Naturally, the `stack()` method provides the opposite operation:

```
In [9]: pop_df.stack()
        # Stack the prescribed level(default -1) from columns to index.
```

```
Out[9]: California 2000    33871648
          2010    37253956
        New York  2000    18976457
          2010    19378102
        Texas    2000    20851820
          2010    25145561
        dtype: int64
```

Seeing this, you might wonder why would we bother with hierarchical indexing at all. The reason is simple: just as we were able to use multi-indexing to represent two-dimensional data within a one-dimensional `Series`, we can also use it to represent data of three or more dimensions in a `Series` or `DataFrame`.

Each extra level in a multi-index represents an extra dimension of data; taking advantage of this property gives us much more flexibility in the types of data we can represent. Concretely, we might want to add another column of demographic data for each state at each year (say, population under 18); with a `MultiIndex` this is as easy as adding another column to the `DataFrame`:

```
In [10]: pop_df = pd.DataFrame({'total': pop,
                                'under18': [9267089, 9284094,
                                              4687374, 4318033,
                                              5906301, 6879014]})

pop_df
```

```
Out[10]:
```

|                   |             | total    | under18 |
|-------------------|-------------|----------|---------|
| <b>California</b> | <b>2000</b> | 33871648 | 9267089 |
|                   | <b>2010</b> | 37253956 | 9284094 |
| <b>New York</b>   | <b>2000</b> | 18976457 | 4687374 |
|                   | <b>2010</b> | 19378102 | 4318033 |
| <b>Texas</b>      | <b>2000</b> | 20851820 | 5906301 |
|                   | <b>2010</b> | 25145561 | 6879014 |

In addition, all the ufuncs and other functionality discussed in previous lecture work with hierarchical indices as well. Here we compute the fraction of people under 18 by year, given the above data:

```
In [11]: f_u18 = pop_df['under18'] / pop_df['total']
f_u18
```

```
Out[11]: California    2000    0.273594
                       2010    0.249211
          New York    2000    0.247010
                       2010    0.222831
          Texas       2000    0.283251
                       2010    0.273568
          dtype: float64
```

```
In [12]: f_u18.unstack()
```

```
Out[12]:
```

|                   | 2000     | 2010     |
|-------------------|----------|----------|
| <b>California</b> | 0.273594 | 0.249211 |
| <b>New York</b>   | 0.247010 | 0.222831 |
| <b>Texas</b>      | 0.283251 | 0.273568 |

This allows us to easily and quickly manipulate and explore even high-dimensional data.

## 2 Methods of MultiIndex Creation

The most straightforward way to construct a multiply indexed `Series` or `DataFrame` is to simply pass a list of two or more index arrays to the constructor. For example:

```
In [13]: df = pd.DataFrame(np.random.rand(4, 2),
                           index=[['a', 'a', 'b', 'b'], [1, 2, 1, 2]],
                           columns=['data1', 'data2'])
df
```

```
Out[13]:
```

|          |          | data1    | data2    |
|----------|----------|----------|----------|
| <b>a</b> | <b>1</b> | 0.886890 | 0.892365 |
|          | <b>2</b> | 0.713180 | 0.289520 |
| <b>b</b> | <b>1</b> | 0.576002 | 0.968222 |
|          | <b>2</b> | 0.224293 | 0.688291 |

```
In [14]: df.index
```

```
Out[14]: MultiIndex([('a', 1),
                     ('a', 2),
                     ('b', 1),
                     ('b', 2)],
                    )
```

The work of creating the `MultiIndex` is done in the background.

Similarly, if you pass a dictionary with appropriate tuples as keys, Pandas will automatically recognize this and use a `MultiIndex` by default:

```
In [15]: data = {('California', 2000): 33871648,
                  ('California', 2010): 37253956,
                  ('Texas', 2000): 20851820,
                  ('Texas', 2010): 25145561,
                  ('New York', 2000): 18976457,
                  ('New York', 2010): 19378102}
se = pd.Series(data)
se
```

```
Out[15]: California 2000    33871648
          2010    37253956
          Texas    2000    20851820
          2010    25145561
          New York 2000    18976457
          2010    19378102
dtype: int64
```

```
In [16]: se.values
```

```
Out[16]: array([33871648, 37253956, 20851820, 25145561, 18976457, 19378102])
```

```
In [17]: se.index
```

```
Out[17]: MultiIndex([('California', 2000),
                     ('California', 2010),
                     ('Texas', 2000),
                     ('Texas', 2010),
                     ('New York', 2000),
                     ('New York', 2010)],
                    )
```

Nevertheless, it is sometimes useful to explicitly create a `MultiIndex` ; we'll see a couple of these methods here.

## 2.1 Explicit MultiIndex constructors

For more flexibility in how the index is constructed, you can instead use the class method constructors available in the `pd.MultiIndex`. For example, as we did before, you can construct the `MultiIndex` from a simple list of arrays giving the index values within each level:

```
In [18]: pd.MultiIndex.from_arrays([['a', 'a', 'b', 'b'], [1, 2, 1, 2]])
```

```
Out[18]: MultiIndex([('a', 1),
                    ('a', 2),
                    ('b', 1),
                    ('b', 2)],
                  )
```

You can construct it from a list of tuples giving the multiple index values of each point:

```
In [19]: pd.MultiIndex.from_tuples([('a', 1), ('a', 2), ('b', 1), ('b', 2)])
```

```
Out[19]: MultiIndex([('a', 1),
                    ('a', 2),
                    ('b', 1),
                    ('b', 2)],
                  )
```

You can even construct it from a Cartesian product of single indices:

```
In [20]: pd.MultiIndex.from_product([['a', 'b'], [1, 2]])
```

```
Out[20]: MultiIndex([('a', 1),
                    ('a', 2),
                    ('b', 1),
                    ('b', 2)],
                  )
```

Any of these objects can be passed as the `index` argument when creating a `Series` or `DataFrame`, or be passed to the `reindex` method of an existing `Series` or `DataFrame`.

## 2.2 MultiIndex level names

Sometimes it is convenient to name the levels of the `MultiIndex`. This can be accomplished by passing the `names` argument to any of the above `MultiIndex` constructors, or by setting the `names` attribute of the index after the fact:

```
In [23]: pop
```

```
Out[23]: California 2000    33871648
          2010    37253956
          New York  2000    18976457
          2010    19378102
          Texas    2000    20851820
          2010    25145561
          dtype: int64
```

```
In [24]: pop.index.names = ['state', 'year']
pop
```

```
Out[24]: state      year
California 2000      33871648
           2010      37253956
New York   2000      18976457
           2010      19378102
Texas      2000      20851820
           2010      25145561
dtype: int64
```

With more involved datasets, this can be a useful way to keep track of the meaning of various index values.

## 2.3 MultiIndex for columns

In a `DataFrame`, the rows and columns are completely symmetric, and just as the rows can have multiple levels of indices, the columns can have multiple levels as well. Consider the following, which is a mock-up of some (somewhat realistic) medical data:

```
In [25]: # hierarchical indices and columns
index = pd.MultiIndex.from_product([[2013, 2014], [1, 2]],
                                   names=['year', 'visit'])
columns = pd.MultiIndex.from_product([['Bob', 'Guido', 'Sue'], ['HR', 'Temp']],
                                    names=['subject', 'type'])

# mock some data
data = np.round(np.random.randn(4, 6), 1)
#round the array for the given number of decimals
data[:, :2] *= 10
data += 37

# create the DataFrame
health_data = pd.DataFrame(data, index=index, columns=columns)
health_data
```

```
Out[25]:
```

|      |       | subject Bob |      | Guido |      | Sue  |      |
|------|-------|-------------|------|-------|------|------|------|
|      | type  | HR          | Temp | HR    | Temp | HR   | Temp |
| year | visit |             |      |       |      |      |      |
| 2013 | 1     | 37.0        | 37.1 | 30.0  | 39.0 | 29.0 | 38.6 |
|      | 2     | 47.0        | 34.8 | 56.0  | 36.5 | 28.0 | 35.8 |
| 2014 | 1     | 18.0        | 37.1 | 41.0  | 37.9 | 30.0 | 36.8 |
|      | 2     | 22.0        | 37.3 | 54.0  | 37.0 | 28.0 | 34.9 |

Here we see where the multi-indexing for both rows and columns can come in *very* handy. This is fundamentally four-dimensional data, where the dimensions are the *subject*, the *measurement type*, the *year*, and the *visit number*. With this in place we can, for example, index the top-level column by the person's name and get a full `DataFrame` containing just that person's information:

```
In [26]: health_data['Bob']
```

```
Out[26]:
```

|  |      | type  | HR   | Temp |
|--|------|-------|------|------|
|  | year | visit |      |      |
|  | 2013 | 1     | 37.0 | 37.1 |
|  |      | 2     | 47.0 | 34.8 |
|  | 2014 | 1     | 18.0 | 37.1 |
|  |      | 2     | 22.0 | 37.3 |

For complicated records containing multiple labeled measurements across multiple times for many subjects (people, countries, cities, etc.) use of hierarchical rows and columns can be extremely convenient!

## 3 Indexing and Slicing a MultiIndex

Indexing and slicing on a `MultiIndex` is designed to be intuitive, and it helps if you think about the indices as added dimensions. We'll first look at indexing multiply indexed `Series`, and then multiply-indexed `DataFrame` s.

### 3.1 Multiply indexed Series

Consider the multiply indexed `Series` of state populations we saw earlier:

```
In [27]: pop
```

```
Out[27]:
```

| state      | year |          |
|------------|------|----------|
| California | 2000 | 33871648 |
|            | 2010 | 37253956 |
| New York   | 2000 | 18976457 |
|            | 2010 | 19378102 |
| Texas      | 2000 | 20851820 |
|            | 2010 | 25145561 |

dtype: int64

We can access single elements by indexing with multiple terms:

```
In [28]: pop['California', 2000]
```

```
Out[28]: 33871648
```

The `MultiIndex` also supports *partial indexing*, or indexing just one of the levels in the index. The result is another `Series`, with the lower-level indices maintained:

```
In [29]: pop['California']
```

```
Out[29]:
```

| year |          |
|------|----------|
| 2000 | 33871648 |
| 2010 | 37253956 |

dtype: int64

Partial slicing is available as well, as long as the `MultiIndex` is sorted :



```
In [30]: pop['California':'New York']
```

```
Out[30]: state      year
California 2000      33871648
           2010      37253956
New York   2000      18976457
           2010      19378102
dtype: int64
```

Partial indexing can be performed on lower levels by passing an empty slice in the first index:

```
In [31]: pop[:, 2000]
```

```
Out[31]: state
California 33871648
New York   18976457
Texas      20851820
dtype: int64
```

```
In [32]: data = {('California', 2000): 33871648,
                  ('California', 2010): 37253956,
                  ('Texas', 2010): 20851820,
                  ('Texas', 2000): 25145561,
                  ('New York', 2000): 18976457,
                  ('New York', 2010): 19378102}
nSort = pd.Series(data)
nSort
```

```
Out[32]: California 2000      33871648
           2010      37253956
Texas          2010      20851820
           2000      25145561
New York       2000      18976457
           2010      19378102
dtype: int64
```

```
In [34]: # nSort['Texas':'New York']
```

Other types of indexing and selection work as well; for example, selection based on Boolean masks:

```
In [35]: pop[pop > 22000000]
```

```
Out[35]: state      year
California 2000      33871648
           2010      37253956
Texas      2010      25145561
dtype: int64
```

Selection based on fancy indexing also works:

```
In [36]: pop[['California', 'Texas']]
```

```
Out[36]: state      year
California 2000      33871648
           2010      37253956
Texas      2000      20851820
           2010      25145561
dtype: int64
```

## 3.2 Multiply indexed DataFrames

A multiply indexed `DataFrame` behaves in a similar manner. Consider our toy medical `DataFrame` from before:

```
In [37]: health_data
```

```
Out[37]:
```

| subject |       | Bob  |      | Guido |      | Sue  |      |
|---------|-------|------|------|-------|------|------|------|
| type    |       | HR   | Temp | HR    | Temp | HR   | Temp |
| year    | visit |      |      |       |      |      |      |
| 2013    | 1     | 37.0 | 37.1 | 30.0  | 39.0 | 29.0 | 38.6 |
|         | 2     | 47.0 | 34.8 | 56.0  | 36.5 | 28.0 | 35.8 |
| 2014    | 1     | 18.0 | 37.1 | 41.0  | 37.9 | 30.0 | 36.8 |
|         | 2     | 22.0 | 37.3 | 54.0  | 37.0 | 28.0 | 34.9 |

Remember that columns are primary in a `DataFrame`, and the syntax used for multiply indexed `Series` applies to the columns. For example, we can recover Guido's heart rate data with a simple operation:

```
In [38]: health_data['Guido', 'HR']
```

```
Out[38]:
```

| year | visit |      |
|------|-------|------|
| 2013 | 1     | 30.0 |
|      | 2     | 56.0 |
| 2014 | 1     | 41.0 |
|      | 2     | 54.0 |

Name: (Guido, HR), dtype: float64

Also, as with the single-index case, we can use the `loc` and `iloc` indexers introduced. For example:

```
In [39]: health_data.iloc[:2, :2]
```

```
Out[39]:
```

| subject |       | Bob  |      |
|---------|-------|------|------|
| type    |       | HR   | Temp |
| year    | visit |      |      |
| 2013    | 1     | 37.0 | 37.1 |
|         | 2     | 47.0 | 34.8 |

These indexers provide an array-like view of the underlying two-dimensional data, but each individual index in `loc` or `iloc` can be passed a tuple of multiple indices. For example:

```
In [40]: health_data.columns
```

```
Out[40]: MultiIndex([( 'Bob',   'HR'),
                    ( 'Bob',   'Temp'),
                    ('Guido',   'HR'),
                    ('Guido',   'Temp'),
                    ( 'Sue',   'HR'),
                    ( 'Sue',   'Temp')],
                    names=['subject', 'type'])
```

```
In [41]: health_data.loc[:, ('Bob', 'HR')]
```

```
Out[41]: year  visit
2013    1      37.0
        2      47.0
2014    1      18.0
        2      22.0
Name: (Bob, HR), dtype: float64
```

Working with slices within these index tuples is not especially convenient; trying to create a slice within a tuple will lead to a syntax error:

You could get around this by building the desired slice explicitly using Python's built-in `slice()` function:

```
In [42]: health_data.loc[(slice(None), 1), (slice(None), 'Temp')]
```

```
Out[42]:
```

|      |       | subject | Bob  | Guido | Sue  |
|------|-------|---------|------|-------|------|
|      |       | type    | Temp | Temp  | Temp |
| year | visit |         |      |       |      |
| 2013 | 1     |         | 37.1 | 39.0  | 38.6 |
| 2014 | 1     |         | 37.1 | 37.9  | 36.8 |

But a better way in this context is to use an `IndexSlice` object, which Pandas provides for precisely this situation. For example:

```
In [43]: idx = pd.IndexSlice
health_data.loc[idx[:, 1], idx[:, 'Temp']]
```

```
Out[43]:
```

|      |       | subject | Bob  | Guido | Sue  |
|------|-------|---------|------|-------|------|
|      |       | type    | Temp | Temp  | Temp |
| year | visit |         |      |       |      |
| 2013 | 1     |         | 37.1 | 39.0  | 38.6 |
| 2014 | 1     |         | 37.1 | 37.9  | 36.8 |

There are so many ways to interact with data in multiply indexed `Series` and `DataFrame`s, and as with many tools in this book the best way to become familiar with them is to try them out!

## 4 Rearranging Multi-Indices

One of the keys to working with multiply indexed data is knowing how to effectively transform the data. There are a number of operations that will preserve all the information in the dataset, but rearrange it for the purposes of various computations. We saw a brief example of this in the `stack()` and `unstack()` methods, but there are many more ways to finely control the rearrangement of data between hierarchical indices and columns, and we'll explore them here.

## 4.1 Sorted and unsorted indices

Earlier, we briefly mentioned a caveat, but we should emphasize it more here. *Many of the* `MultiIndex` *slicing operations will fail if the index is not sorted.* Let's take a look at this here.

We'll start by creating some simple multiply indexed data where the indices are *not lexicographically sorted*:

```
In [44]: index = pd.MultiIndex.from_product(['a', 'c', 'b'], [2, 1])
data = pd.Series(np.random.rand(6), index=index)
data.index.names = ['char', 'int']
data
```

```
Out[44]: char  int
a         2    0.947427
          1    0.884243
c         2    0.817343
          1    0.498541
b         2    0.041736
          1    0.016117
dtype: float64
```

If we try to take a partial slice of this index, it will result in an error:

```
In [45]: try:
          data['a':'b']
        except KeyError as e:
            print(type(e))
            print(e)
```

```
<class 'pandas.errors.UnsortedIndexError'>
'Key length (1) was greater than MultiIndex lexsort depth (0)'
```

Although it is not entirely clear from the error message, this is the result of the `MultiIndex` not being sorted. For various reasons, partial slices and other similar operations require the levels in the `MultiIndex` to be in sorted (i.e., lexicographical) order. Pandas provides a number of convenience routines to perform this type of sorting; examples are the `sort_index()` and `sortlevel()` methods of the `DataFrame`. We'll use the simplest, `sort_index()`, here:

```
In [46]: data = data.sort_index()
data
```

```
Out[46]: char  int
a         1    0.884243
          2    0.947427
b         1    0.016117
          2    0.041736
c         1    0.498541
          2    0.817343
dtype: float64
```

With the index sorted in this way, partial slicing will work as expected:

```
In [47]: data['a':'b']
```

```
Out[47]: char  int
a         1    0.884243
          2    0.947427
b         1    0.016117
          2    0.041736
dtype: float64
```

## 4.2 Stacking and unstacking indices

As we saw briefly before, it is possible to convert a dataset from a stacked multi-index to a simple two-dimensional representation, optionally specifying the level to use:

```
In [48]: pop
```

```
Out[48]: state      year
California 2000      33871648
           2010      37253956
New York   2000      18976457
           2010      19378102
Texas      2000      20851820
           2010      25145561
dtype: int64
```

```
In [49]: pop.unstack(level=0)
```

```
Out[49]:
```

|      | state | California | New York | Texas    |
|------|-------|------------|----------|----------|
| year |       |            |          |          |
| 2000 |       | 33871648   | 18976457 | 20851820 |
| 2010 |       | 37253956   | 19378102 | 25145561 |

```
In [50]: pop.unstack(level=1)
```

```
Out[50]:
```

| year       | 2000     | 2010     |
|------------|----------|----------|
| state      |          |          |
| California | 33871648 | 37253956 |
| New York   | 18976457 | 19378102 |
| Texas      | 20851820 | 25145561 |

The opposite of `unstack()` is `stack()`, which here can be used to recover the original series:

```
In [51]: pop.unstack().stack()
```

```
Out[51]: state      year
California 2000      33871648
           2010      37253956
New York   2000      18976457
           2010      19378102
Texas      2000      20851820
           2010      25145561
dtype: int64
```

## 4.3 Index setting and resetting

Another way to rearrange hierarchical data is to **turn the index labels into columns**; this can be accomplished with the `reset_index` method. Calling this on the population dictionary will result in a `DataFrame` with a *state* and *year* column holding the information that was formerly in the index.

In [52]:

```
pop
```

Out[52]:

```
state      year
California 2000    33871648
           2010    37253956
New York   2000    18976457
           2010    19378102
Texas      2000    20851820
           2010    25145561
dtype: int64
```

In [53]:

```
pop.reset_index()
```

Out[53]:

|   | state      | year | 0        |
|---|------------|------|----------|
| 0 | California | 2000 | 33871648 |
| 1 | California | 2010 | 37253956 |
| 2 | New York   | 2000 | 18976457 |
| 3 | New York   | 2010 | 19378102 |
| 4 | Texas      | 2000 | 20851820 |
| 5 | Texas      | 2010 | 25145561 |

For clarity, we can optionally specify the name of the data for the column representation:

In [54]:

```
pop_flat = pop.reset_index(name='population')
pop_flat
```

Out[54]:

|   | state      | year | population |
|---|------------|------|------------|
| 0 | California | 2000 | 33871648   |
| 1 | California | 2010 | 37253956   |
| 2 | New York   | 2000 | 18976457   |
| 3 | New York   | 2010 | 19378102   |
| 4 | Texas      | 2000 | 20851820   |
| 5 | Texas      | 2010 | 25145561   |

Often when working with data in the real world, the raw input data looks like this and it's useful to build a `MultiIndex` from the column values. This can be done with the `set_index` method of the `DataFrame`, which returns a multiply indexed `DataFrame`:

In [55]:

```
# pop_flat.set_index Set the DataFrame index (row labels) using one or more existing columns.
pop_flat.set_index(['state', 'year'])
```

Out[55]:

|  |            | population |          |
|--|------------|------------|----------|
|  | state      | year       |          |
|  | California | 2000       | 33871648 |
|  |            | 2010       | 37253956 |
|  | New York   | 2000       | 18976457 |
|  |            | 2010       | 19378102 |
|  | Texas      | 2000       | 20851820 |
|  |            | 2010       | 25145561 |

## 5 Data Aggregations on Multi-Indices

We've previously seen that Pandas has built-in data aggregation methods, such as `mean()`, `sum()`, and `max()`. For hierarchically indexed data, these can be passed a `level` parameter that controls which subset of the data the aggregate is computed on.

For example, let's return to our health data:

```
In [56]: health_data
```

```
Out[56]:
```

|      |       | subject Bob |      | Guido |      | Sue  |      |
|------|-------|-------------|------|-------|------|------|------|
|      | type  | HR          | Temp | HR    | Temp | HR   | Temp |
| year | visit |             |      |       |      |      |      |
| 2013 | 1     | 37.0        | 37.1 | 30.0  | 39.0 | 29.0 | 38.6 |
|      | 2     | 47.0        | 34.8 | 56.0  | 36.5 | 28.0 | 35.8 |
| 2014 | 1     | 18.0        | 37.1 | 41.0  | 37.9 | 30.0 | 36.8 |
|      | 2     | 22.0        | 37.3 | 54.0  | 37.0 | 28.0 | 34.9 |

Perhaps we'd like to average-out the measurements in the two visits each year. We can do this by naming the index level we'd like to explore, in this case the year:

```
In [57]: data_mean = health_data.mean(level='year')
data_mean
```

```
Out[57]:
```

|      | subject Bob |       | Guido |       | Sue  |       |
|------|-------------|-------|-------|-------|------|-------|
| type | HR          | Temp  | HR    | Temp  | HR   | Temp  |
| year |             |       |       |       |      |       |
| 2013 | 42.0        | 35.95 | 43.0  | 37.75 | 28.5 | 37.20 |
| 2014 | 20.0        | 37.20 | 47.5  | 37.45 | 29.0 | 35.85 |

By further making use of the `axis` keyword, we can take the mean among levels on the columns as well:

```
In [58]: data_mean.mean(axis=1, level='type')
```

```
Out[58]:
```

| type | HR        | Temp      |
|------|-----------|-----------|
| year |           |           |
| 2013 | 37.833333 | 36.966667 |
| 2014 | 32.166667 | 36.833333 |

Thus in two lines, we've been able to find the average heart rate and temperature measured among all subjects in all visits each year.